

## 1. Data wrangling using the tidyverse (classwork problems)

Load the nycflights13 library (will have to install the nycflights13 package first) which contains flight arrival and departure data in a table called `flights`. Apply the tidyverse's data wrangling verbs to answer these questions. For each question, **give only the (one line) code**.

- a. List data only for flights that departed on February 12,

**Commands:**

```
> flights %>% filter( month==2, day==12)
```

**Screen Shot:**

```
> flights %>% filter( month==2, day==12)
# A tibble: 893 x 19
  year month   day dep_time sched_dep_time dep_delay arr_time sched_arr_time arr_delay carrier
  <int> <int> <int>   <int>         <int>         <dbl>   <int>         <int>         <dbl> <chr>
1  2013     2    12      17             2245          92     122             2356          86 B6
2  2013     2    12     506             500           6     703             648          15 US
3  2013     2    12     520             525         -5     837             820          17 UA
4  2013     2    12     524             530         -6     922             831          51 UA
5  2013     2    12     535             540         -5     950             1016         -26 B6
6  2013     2    12     539             540         -1     828             850         -22 AA
7  2013     2    12     551             600         -9     645             708         -23 B6
8  2013     2    12     552             600         -8     925             910          15 AA
9  2013     2    12     553             600         -7     652             703         -11 US
10 2013     2    12     555             600         -5     903             911          -8 B6
# ... with 883 more rows, and 9 more variables: flight <int>, tailnum <chr>, origin <chr>,
#   dest <chr>, air_time <dbl>, distance <dbl>, hour <dbl>, minute <dbl>, time_hour <dtm>
> |
```

- b. List data only for flights that were delayed (both arrival and departure) by more than 2 hours.

**Commands:**

```
> flights %>% filter(dep_delay >120, arr_delay>120)
```

**Screen Shot:**

```
> flights %>% filter(dep_delay >120, arr_delay>120)
# A tibble: 8,335 x 19
  year month   day dep_time sched_dep_time dep_delay arr_time sched_arr_time arr_delay carrier
  <int> <int> <int> <int>         <int>         <dbl> <int>         <int>         <dbl> <chr>
1  2013     1     1     848           1835         853    1001           1950         851 MQ
2  2013     1     1     957           733          144    1056           853          123 UA
3  2013     1     1    1114           900          134    1447           1222          145 UA
4  2013     1     1    1815          1325          290    2120           1542          338 EV
5  2013     1     1    1842          1422          260    1958           1535          263 EV
6  2013     1     1    1856          1645          131    2212           2005          127 AA
7  2013     1     1    1934          1725          129    2126           1855          151 MQ
8  2013     1     1    1938          1703          155    2109           1823          166 EV
9  2013     1     1    1942          1705          157    2124           1830          174 MQ
10 2013     1     1    2006          1630          216    2230           1848          222 EV
# ... with 8,325 more rows, and 9 more variables: flight <int>, tailnum <chr>, origin <chr>,
#   dest <chr>, air_time <dbl>, distance <dbl>, hour <dbl>, minute <dbl>, time_hour <dtm>
> |
```

- c. List data only for flights that were delayed (either arrival or departure) by more than 2 hours.

**Commands:**

```
> flights %>% filter(dep_delay>120 | arr_delay>120)
```

**Screen Shot:**

```
> flights %>% filter(dep_delay>120 | arr_delay>120)
# A tibble: 11,422 x 19
  year month   day dep_time sched_dep_time dep_delay arr_time sched_arr_time arr_delay carrier
  <int> <int> <int> <int>         <int>         <dbl> <int>         <int>         <dbl> <chr>
1  2013     1     1     811           630          101    1047           830          137 MQ
2  2013     1     1     848           1835         853    1001           1950         851 MQ
3  2013     1     1     957           733          144    1056           853          123 UA
4  2013     1     1    1114           900          134    1447           1222          145 UA
5  2013     1     1    1505          1310          115    1638           1431          127 EV
6  2013     1     1    1525          1340          105    1831           1626          125 B6
7  2013     1     1    1540          1338          122    2020           1825          115 B6
8  2013     1     1    1549          1445           64    1912           1656          136 EV
9  2013     1     1    1558          1359          119    1718           1515          123 EV
10 2013     1     1    1732          1630           62    2028           1825          123 EV
# ... with 11,412 more rows, and 9 more variables: flight <int>, tailnum <chr>, origin <chr>,
#   dest <chr>, air_time <dbl>, distance <dbl>, hour <dbl>, minute <dbl>, time_hour <dtm>
> |
```

- d. List data only for flights that were operated by United, American, or Delta.

**Commands:**

```
> flights %>% filter(carrier=="UA" | carrier=="AA" | carrier=="DL")
```

**OR**

```
> flights %>% filter(carrier %in% c("UA", "AA", "DL"))
```

### Screen Shot:

```
> flights %>% filter(carrier=="UA" | carrier=="AA" | carrier=="DL")
# A tibble: 139,504 x 19
  year month   day dep_time sched_dep_time dep_delay arr_time sched_arr_time arr_delay
  <int> <int> <int>   <int>         <int>         <dbl>   <int>         <int>         <dbl>
1  2013     1     1     517             515           2     830             819           11
2  2013     1     1     533             529           4     850             830           20
3  2013     1     1     542             540           2     923             850           33
4  2013     1     1     554             600          -6     812             837          -25
5  2013     1     1     554             558          -4     740             728           12
6  2013     1     1     558             600          -2     753             745           8
7  2013     1     1     558             600          -2     924             917           7
8  2013     1     1     558             600          -2     923             937          -14
9  2013     1     1     559             600          -1     941             910           31
10 2013     1     1     559             600          -1     854             902           -8
# ... with 139,494 more rows, and 10 more variables: carrier <chr>, flight <int>,
#   tailnum <chr>, origin <chr>, dest <chr>, air_time <dbl>, distance <dbl>,
#   hour <dbl>, minute <dbl>, time_hour <dtm>
> |
```

- e. Sort data in order of fastest flights.

### Commands:

```
> flights %>% arrange(arr_time)
```

### Screen Shot:

```
> flights %>% arrange(arr_time)
# A tibble: 336,776 x 19
  year month   day dep_time sched_dep_time dep_delay arr_time sched_arr_time arr_delay
  <int> <int> <int>   <int>         <int>         <dbl>   <int>         <int>         <dbl>
1  2013     1     2     2130             2130           0         1             18          -17
2  2013     1    11     2157             2000        117         1             2208        113
3  2013     1    11     2253             2249           4         1             2357           4
4  2013     1    14     2122             2130          -8         1              2          -1
5  2013     1    14     2246             2250          -4         1              7          -6
6  2013     1    15     2304             2245         19         1             2357           4
7  2013     1    16     2018             2025          -7         1             2329          32
8  2013     1    16     2303             2245         18         1             2357           4
9  2013     1    19     2107             2110          -3         1             2355           6
10 2013     1    22     2246             2249          -3         1             2357           4
# ... with 336,766 more rows, and 10 more variables: carrier <chr>, flight <int>,
#   tailnum <chr>, origin <chr>, dest <chr>, air_time <dbl>, distance <dbl>,
#   hour <dbl>, minute <dbl>, time_hour <dtm>
> |
```

- f. Sort data in order of longest flights.

### Commands:

```
> flights %>% arrange(desc(arr_time))
```

### Screen Shot:

```
> flights %>% arrange(desc(arr_time))
# A tibble: 336,776 x 19
  year month   day dep_time sched_dep_time dep_delay arr_time sched_arr_time arr_delay
  <int> <int> <int>   <int>         <int>         <dbl>   <int>         <int>         <dbl>
1  2013     1     1    2209           2155         14    2400           2337          23
2  2013     1     5    2116           2130        -14    2400            18        -18
3  2013     1    13    2243           2129         74    2400           2224          96
4  2013     1    16    2138           2107         31    2400           2322          38
5  2013     1    17    2256           2249          7    2400           2357           3
6  2013     1    22    2212           2055         77    2400           2250          70
7  2013     1    22    2249           2125         84    2400           2250          70
8  2013     1    25    2055           1725        210    2400           1933         267
9  2013     1    28    2303           2250         13    2400           2354           6
10 2013     1    30    2155           1915        160    2400           2137         143
# ... with 336,766 more rows, and 10 more variables: carrier <chr>, flight <int>,
#   tailnum <chr>, origin <chr>, dest <chr>, air_time <dbl>, distance <dbl>,
#   hour <dbl>, minute <dbl>, time_hour <dtm>
>
```

- g. Show only the origin and destination of flights sorted by longest flights.

**Commands:**

```
> flights %>% arrange(desc(air_time)) %>% select(origin, dest)
```

**Screen Shot:**

```
> flights %>% arrange(desc(air_time)) %>% select(origin, dest)
# A tibble: 336,776 x 2
  origin dest
  <chr>   <chr>
1 EWR     HNL
2 JFK     HNL
3 JFK     HNL
4 JFK     HNL
5 JFK     HNL
6 JFK     HNL
7 EWR     HNL
8 JFK     HNL
9 JFK     HNL
10 EWR     HNL
# ... with 336,766 more rows
> |
```

- h. Add a new variable that indicates the total delay (both departure and arrival delay).

**Commands:**

```
> flights %>% mutate( total_delay = dep_delay+arr_delay)
```

### Screen Shot:

```
> flights %>% mutate( total_delay = dep_delay+arr_delay)
# A tibble: 336,776 x 20
   year month   day dep_time sched_dep_time dep_delay arr_time sched_arr_time arr_delay
   <int> <int> <int>   <int>         <int>         <dbl>   <int>         <int>         <dbl>
1  2013     1     1     517           515           2     830           819           11
2  2013     1     1     533           529           4     850           830           20
3  2013     1     1     542           540           2     923           850           33
4  2013     1     1     544           545          -1    1004          1022          -18
5  2013     1     1     554           600          -6     812           837          -25
6  2013     1     1     554           558          -4     740           728           12
7  2013     1     1     555           600          -5     913           854           19
8  2013     1     1     557           600          -3     709           723          -14
9  2013     1     1     557           600          -3     838           846           -8
10 2013     1     1     558           600          -2     753           745            8
# ... with 336,766 more rows, and 11 more variables: carrier <chr>, flight <int>,
#   tailnum <chr>, origin <chr>, dest <chr>, air_time <dbl>, distance <dbl>,
#   hour <dbl>, minute <dbl>, time_hour <dtm>, total_delay <dbl>
>
```

- i. Show only the origin and destination of flights sorted by descending order of total delay.

### Commands:

```
> flights %>% mutate(total_delay=dep_delay+arr_delay) %>%
  arrange(desc(total_delay)) %>% select (origin,dest)
```

### Screen Shot:

```
> flights %>% mutate(total_delay=dep_delay+arr_delay) %>% arrange(desc(total_delay)) %>% select (origin,dest)
# A tibble: 336,776 x 2
   origin dest
   <chr>   <chr>
1 JFK     HNL
2 JFK     CMH
3 EWR     ORD
4 JFK     SFO
5 JFK     CVG
6 JFK     TPA
7 LGA     MSP
8 LGA     ATL
9 EWR     MIA
10 EWR     ORD
# ... with 336,766 more rows
>
```

- j. Show only the origin and destination of 10 most delayed flights (Hint: use the min\_rank() function which assigns ranks 1, 2, 3, ...).

### Commands:

```
> flights %>% mutate(total_delay = arr_delay + dep_delay) %>%  
mutate(total_delay_rank = min_rank(desc(total_delay))) %>%  
arrange(total_delay_rank) %>% select(origin, dest) %>% slice(1:10) %>% view()
```

### Screen Shot:

```
> flights %>% mutate(total_delay = arr_delay + dep_delay) %>% mutate(total_delay_rank = min_rank(desc(total_delay))) %>% arrange(to  
tal_delay_rank) %>% select(origin, dest) %>% slice(1:10)  
# A tibble: 10 x 2  
  origin dest  
  <chr> <chr>  
1 JFK   HNL  
2 JFK   CMH  
3 EWR   ORD  
4 JFK   SFO  
5 JFK   CVG  
6 JFK   TPA  
7 LGA   MSP  
8 LGA   ATL  
9 EWR   MIA  
10 EWR  ORD  
> |
```

- k. Show the average total delay (excluding NA values) for every origin city.

### Commands:

```
> flights %>% na.exclude() %>% mutate(total_delay = arr_delay + dep_delay) %>%  
group_by(origin) %>% summarise(total_delay_mean = mean(total_delay))
```

### Screen Shot:

```
> flights %>% na.exclude() %>% mutate(total_delay = arr_delay + dep_delay) %>% group_by(origin) %>% summarise(total_delay_mean = me  
an(total_delay))  
# A tibble: 3 x 2  
  origin total_delay_mean  
  <chr> <dbl>  
1 EWR   24.1  
2 JFK   17.6  
3 LGA   16.1  
>
```

- l. Show the average total delay (excluding NA values) for every origin-destination city pair.

### Commands:

```
> flights %>% na.exclude() %>% mutate(total_delay = arr_delay + dep_delay) %>%  
group_by(origin, dest) %>% summarise(total_delay_mean = mean(total_delay))
```

### Screen Shot:



```

> flights %>% na.exclude() %>% mutate(total_delay = arr_delay + dep_delay) %>% group_by(origin, dest) %>% summarise(total_delay_mean = mean(total_delay))
# A tibble: 223 x 3
# Groups:   origin [?]
  origin dest total_delay_mean
  <chr>   <chr>             <dbl>
1 EWR    ALB                37.8
2 EWR    ANC                10.4
3 EWR    ATL                28.6
4 EWR    AUS                 11
5 EWR    AVL                17.4
6 EWR    BDL                24.8
7 EWR    BNA                30.3
8 EWR    BOS                17.3
9 EWR    BQN                34.5
10 EWR   BTV                30.0
# ... with 213 more rows
> |

```

## 2. Data reshaping using the tidyverse

- a. Consider the attached .csv file “horse\_racing.csv” which contains data related to horse racing licensing in New York<sup>1</sup>. The `License` column has two types of values: license numbers and receipt numbers. Load the dataset and transform it such that this column is split into two:
  - i. `LicenseOrReceipt`: a factor with two levels “License” and “Receipt”
  - ii. `Number`: numeric column with the license/receipt number

Show (1) your code, and (2) copy & paste the output of the function `str()` on your final table.

```

> hracesep <- separate(hrace, License, into = c("LicenseOrReceipt", "Number"), sep = "#")
> str(hracesep)
'data.frame':   24191 obs. of  8 variables:
 $ PersonID      : int  120384 148737 200788 200514 59203 155736 143125 143125 143125 195645
 ...
 $ Name          : Factor w/ 20487 levels "0MAR MEHIDI",...: 1 2 3 4 5 6 7 7 7 8 ...
 $ Occupation    : Factor w/ 97 levels "APPRENTICE JOCKEY",...: 8 8 78 57 58 58 67 67 67 82 ..
 .
 $ Eligibility   : Factor w/ 2 levels "ABLE TO PARTICIPATE",...: 1 1 1 1 1 1 2 2 2 1 ...
 $ Division      : Factor w/ 2 levels "HARNESS", "THOROUGHBRED": 2 2 1 2 2 2 1 1 1 1 ...
 $ LicenseOrReceipt: chr  "LICENSE " "RECEIPT " "RECEIPT " "RECEIPT " ...
 $ Number        : chr  " 1522818" " 1462171" " 1462094" " 1449814" ...
 $ Expires       : Factor w/ 1134 levels "1/1/2020", "1/1/2021",...: 613 214 363 350 682 1066 1
 91 191 191 910 ...

```

- b. Consider the attached .csv file, “language\_diversity.csv,” which contains data on the diversity of languages in different countries and other parameters<sup>2</sup>.
  - a. Is the data “tidy”? Explain your answer in 2-3 sentences.

<sup>1</sup> Original dataset: <https://data.ny.gov/Government-Finance/Horse-Racing-Licensing/cz9u-yj7m/data>

<sup>2</sup> Dataset from: [https://github.com/jvcasillas/untidydata#language\\_diversity](https://github.com/jvcasillas/untidydata#language_diversity)

It's not tidy because the variables MGS, Population, Stations, Std, langs and area should be made into columns instead of rows. Observations are unique values of MGS, Population, Stations, Std, langs and area in a given country and continent.

- b. Convert the data to tidy data. Show (1) your code, and (2) copy & paste the output of the function `str` on your final table.

**Code:**

```
> lang %>% spread(Measurement, Value)
```

**Output:**

```
> str(lang %>% spread(Measurement, Value))
'data.frame': 74 obs. of 8 variables:
 $ Continent : Factor w/ 4 levels "Africa","Americas",...: 1 1 1 1 1 1 1 1 1 1 ...
 $ Country   : Factor w/ 74 levels "Algeria","Angola",...: 1 2 5 7 9 11 12 13 15 17 ...
 $ Area      : num 2381741 1246700 112622 581730 274000 ...
 $ Langs     : num 18 42 52 27 75 275 94 126 60 75 ...
 $ MGS       : num 6.6 6.22 7.14 4.6 5.17 9.17 8.08 4 9.6 8.67 ...
 $ Population: num 25660 10303 4889 1348 9242 ...
 $ Stations  : num 102 50 7 10 6 35 13 11 10 9 ...
 $ Std       : num 2.29 1.87 0.99 1.69 1.07 1.75 1.21 1.81 1.69 1.25 ...
> |
```

- c. Consider the attached .csv file, “diseases.csv,” which contains data from Australia on hospitalizations<sup>3</sup>.

Diseases	Patientdays_Y2 015-16	Separations_Y 2015-16	Patientdays_Y2 016-17	Separations_Y 2016-17
1 Certain infectious and parasitic diseases (A00-B99)	694,007	170,095	771,770	186,034
2 Neoplasms (C00-D48)	2,223,563	666,594	2,235,045	684,075
3 Diseases of the blood and blood-forming organs and certain disorders involving the immune mechanism (D50-D89)	317,085	175,590	335,699	190,568

The first few rows are shown above. Load this file and convert the table to the tidy format shown below. Note the new column names. Show (1) your code, and (2) copy & paste the output of the function `str` on your final table. (Hint: this will require multiple transforms from `gather/separate/select`. Read the file with `read_csv`, not `read.csv`)

<sup>3</sup> Dataset from:

<https://www.aihw.gov.au/reports/hospitals/principal-diagnosis-data-cubes/contents/data-cubes>



Diseases	Year	Patientdays	Separations
1 Certain infectious and parasitic diseases (A00-B99)	Y2015-16	694,007	170,095
1 Certain infectious and parasitic diseases (A00-B99)	Y2016-17	771,770	186,034
2 Neoplasms (C00-D48)	Y2015-16	2,223,563	666,594
2 Neoplasms (C00-D48)	Y2016-17	2,235,045	684,075

```

> diseasestemp <- select(diseases, c(1,2,4))
> diseasestemp2 <- select(diseases, c(1,3,5))
> fixdiseases1 <- separate(diseasestemp, "Patientdays_Y2015-16",
into=c("Y2015-16","Patientdays"), sep="_")
> fixdiseases1 <- separate(fixdiseases1, "Patientdays_Y2016-17",
into=c("Y2016-17","Patientdays"), sep="_")
> fixdiseases1 <- gather(fixdiseases1, Year, Patientdays, c(2:4))
> fixdiseases1 <- fixdiseases1[-c(43:63),] #deletes extra rows that were created
> fixdiseases2 <- separate(diseasestemp2, "Separations_Y2015-16",
into=c("Y2015-16","Separations"), sep="_")
> fixdiseases2 <- gather(fixdiseases2, Year, Separations, c(2:4))
> fixdiseases2 <- fixdiseases2[-c(43:63),] #deletes extra rows that were created
> mergedDiseases <- merge(fixdiseases1, fixdiseases2, by=c("Diseases", "Year")) #merge the 2
separate datasets into the complete dataset

```

```

> str(mergedDiseases)
'data.frame': 42 obs. of 4 variables:
 $ Diseases : chr "1 Certain infectious and parasitic diseases (A00-B99)" "1 Certain infectious and par
asitic diseases (A00-B99)" "10 Diseases of the respiratory system (J00-J99)" "10 Diseases of the respirato
ry system (J00-J99)" ...
 $ Year : chr "Y2015-16" "Y2016-17" "Y2015-16" "Y2016-17" ...
 $ Patientdays: chr "694007" "771770" "1700645" "1788798" ...
 $ Separations: chr "170095" "186034" "467780" "498853" ...
> |

```